

Highlights of this Solr release include:¶

- Two dimensional routed aliases are now available for organizing collections based on the data values of two fields
- SPLITSHARD implements a new splitByPrefix option that takes into account the actual document distribution when using compositelds
- QueryElevationComponent can have query rules configured with match="subset" wherein the words need only match a subset of the query's words and in any order
- Command line option to export documents to a file
- Support deterministic replica routing preferences for better cache usage
- Ability to query aliases in Solr Admin UI
- JWTAuthPlugin supports multiple JWKS endpoints and multiple IdP issuers
- JSON faceting now supports arbitrary ranges for range facets
- Support integral plots, cosine distance and string truncation with math expressions (Joel Bernstein)
- New cat() stream source to create tuples from lines in local files
- Add upper, lower, trim and split Stream Evaluators
- Add CsvStream, TsvStream Streaming Expressions and supporting Stream Evaluators
- Add CaffeineCache, an efficient implementation of SolrCache
- Live SPLITSHARD can lose updates due to cluster state change between checking if the current shard is active and later checking if there are any sub-shard leaders to forward the update to
- Fix for SPLITSHARD (async) with failures in underlying sub-operations can result in data loss
- Allow dynamic resizing of SolrCache-s
- Allow optional redaction of data saved by 'bin/solr autoscaling - save'
- Optimized large managed schema modifications (internal $O(n^2)$ problem)
- Max idle time support for SolrCache implementations
- Add Prometheus Exporter GC and Heap options
- SSL: Adding Enabling/Disabling client's hostname verification config
- Introducing SolrClient.ping(collection) in SolrJ
- Fix for CDCR bootstrap not replicating index to the replicas of target cluster
- Fixed a race condition when initializing metrics for new security plugins on security.json change
- Fixed JWTAuthPlugin to update metrics prior to continuing w/ other filters or returning error

Fixed distributed grouping when multiple 'fl' params are specified
JMX MBeans are not exposed because of race condition between creating platform mbean server and registering mbeans
Fix for class-cast issues during atomic-update 'remove regex' operations
Fix for multi-node race condition to create/remove nodeLost markers
Fix for too many cascading calls to remote servers, which can bring down nodes
Fix for MOVEREPLICA ignoring replica type and always adding 'nrt' replicas
Fix: DistributedZkUpdateProcessor should propagate URP.finish() lifecycle (regression since 8.1)

6 September 2019 - New mailing lists¶

The Lucene project has added two new announce mailing lists, issues@lucene.apache.org and builds@lucene.apache.org. High-volume automated emails from our bug tracker, JIRA and GitHub will be moved from the [dev@](mailto:dev@lucene.apache.org) list to [issues@](mailto:issues@lucene.apache.org) and automated emails from our Jenkins CI build servers will be moved from the [dev@](mailto:dev@lucene.apache.org) list to [builds@](mailto:builds@lucene.apache.org).

This is an effort to reduce the sometimes overwhelming email volume on our main development mailing list and thus make it easier for the community to follow important discussions by humans on the dev@lucene.apache.org list.

Everyone who wants to continue receiving these automated emails should sign up for one or both of the two new lists. Sign-up instructions can be found on the [Lucene-java](#) and [Solr](#) web sites.

26 July 2019 - Apache Lucene 8.2.0 and Apache Solr 8.2.0 Available¶

The Lucene PMC is pleased to announce the release of Apache Lucene 8.2.0 and Apache Solr 8.2.0

Lucene can be downloaded from <https://lucene.apache.org/core/downloads.html> and Solr can be downloaded from <https://lucene.apache.org/solr/downloads.html>

Highlights of this Lucene release include:¶

API Changes:¶

Intervals queries has been moved from the sandbox to the queries module.

New Features¶

New XYShape Field and Queries for indexing and querying general cartesian geometries.

Snowball stemmer/analyzer for the Estonian language.
Provide a FeatureSortfield to allow sorting search hits by descending value of a feature.
Add new KoreanNumberFilter that can change Hangul character to number and process decimal point.
Add doc-value support to range fields.
Add monitor subproject (previously Luwak monitoring library) that allows a stream of documents to be matched against a set of registered queries in an efficient manner.
Add a numeric range query in sandbox that takes advantage of index sorting. Add a numeric range query in sandbox that takes advantage of index sorting.

Optimizations ¶

Use exponential search instead of binary search in IntArrayDocIdSet#advance method.
Use incoming thread for execution if IndexSearcher has an executor. Now caller threads execute at least one search on an index even if there is an executor provided to minimize thread context switching.
New storing strategy for BKD tree leaves with low cardinality that can lower storage costs and it can be used at search time to speed up queries.
Load frequencies lazily only when needed in BlockDocsEnum and BlockImpactsEverythingEnum.
Phrase queries now leverage impacts.

Highlights of this Solr release include: ¶

New features ¶

Add an update param failOnVersionConflicts=false to updates not fail when there is a version conflict
Add facet2D Streaming Expression.
Preferred replicas on nodes with same system properties as the query master
OpenTracing support for Solr
Raw index data analysis tool (extension of COLSTATUS collection command).
Add recNum Stream Evaluator.
Allow zplot to visualize 2D clusters and convex hulls.
Add a field type for Estonian language to default managed_schema, document about Estonian language analysis in Solr Ref Guide

Bug Fixes ¶

Intermittent 401's for internode requests with basic auth enabled.
In 8.1, Atomic Updates were broken (NPE) when the schema declared the new nest_path field even if you weren't using nested

docs. In-place updates were not affected (worked)

Fix atomic update encoding issue for UUID, enum, bool, and binary fields.

Impossible to delete a collection with the same name as an existing alias. This fixes also a bug in REINDEXCOLLECTION when used with removeSource=true which could lead to a data loss.